


Built at the MTSU Great Robotics Lab. Every robot the lab owns reports into one platform; agents on it harvest, filter, label and train; a third agent watches the other two. This is what runs today, what the robots have recorded, and what six months of running it taught us.


WHAT EACH PROJECT IS, AND WHAT IT IS FOR

Weed Detection



Two agents work the same corpus without being asked. One goes out to public sources, brings weed imagery back and audits what it brought; the other trains a detector on whatever survives the audit and scores it. The registry behind them holds more than 150,000 labelled weed images, and the detector that comes out is small enough to ride on the robots themselves.

241 Robot



A tracked rover built for the row: narrow enough to turn inside one and tough enough to cross the beds between them, so it can reach parts of a field a tractor cannot. It carries a camera, GPS and two inertial units and streams all of them live into the platform. That stream is what an autonomous-driving model for the field is trained on, and a signed-in user drives the rover from a plain browser.

Laser Cart



A twin-laser weeder. A YOLO detector finds the weed in the frame, the two galvanometer lasers are aimed at it and fired, and the weed is killed where it stands -- no chemical, and nothing touches the crop beside it. One camera looks down over the work zone and one looks forward, and both stream into the platform while the cart's own operator drives and fires it.

Humanoid and Quadruped



A humanoid and a quadruped, taken out to the farm to record where wheels do not go. Their project is already open and takes the same ingest contract the rover and the cart use, so the day either of them opens a session it is a governed dataset like any other.

Figure 1. Every project on the platform was opened the same way, sends through the same ingest contract, and is trained and evaluated by the same agents.

HOW A ROUND RUNS, AND WHAT IT DOES WHEN IT IS STUCK




Figure 2. The loop an agent runs to answer a need, drawn as a loop. A need arrives, the agent works out what data would answer it, goes and gets it, trains on what survives the audit and scores it, then reads its own result and decides what to ask for next. When it cannot decide it climbs the ladder inside the loop, each rung reading the one below it and free to overrule it. The outward run is solid because it runs unattended today. The return run is dashed because it is the one edge the platform has not closed: the decision is recorded and nothing reads it back, and closing it is the work this poster is about.

ONE PLATFORM, ANY ROBOT, ANY DATASET

- **agentAI is one platform for every robot the lab owns**
A robot, a drone or a sensor rig is a project. It owns its keys, its people and a locked dataset registry, and the agents harvest, filter, label and train against that registry. Adding a robot is a configuration change.
- **Weed detection is one project here, the humanoid is still empty**
Open a project for any data type and any field, upload, and get an automatic diagnosis before choosing what to do next. Weed detection is one project here; so are the rover and the cart.
- **Adding the next robot takes a project, a key and a session**
A new robot needs a project, an API key and a session start. The ingest contract is the same for every vehicle and the platform does not care what the robot is. Empty projects for a mobile robot and a humanoid already sit on it, waiting for a first upload.
- **Tracked rover and laser cart are the two robots reporting in**
Robot 241 is a tracked rover that streams camera, GPS and two IMUs. The laser cart carries a galvanometer laser, a down camera over the work zone and a forward camera.
- **A drive becomes a dataset the moment it opens**
A vehicle opens a session and the platform registers it as a governed dataset at that moment, while the drive is still running. Frames and sensor batches arrive over one API key. Robot 241 sends camera, GPS and two IMUs; the laser cart sends camera, laser state, vehicle telemetry and its own detections, and no position at all.


A BRAIN THAT GETS STUCK AND ASKS UPWARD



THE AGENTS RAN THE LOOP WITH NOBODY IN THE ROOM

- **Fifteen rounds ran with nobody in the room**
Two cluster jobs, one collecting and one training, ran fifteen rounds against the locked registry with nobody in the room. Every round reported success. This part is built and is switched off today.

MORE DATA STOPPED HELPING, SO WE WENT AFTER THE LABELS



Starting fresh instead of from the last round's weights recovers +0.0287, 5.1 times this recipe's seed spread. That is the change the next campaign makes.

MORE DATA STOPPED HELPING, SO WE WENT AFTER THE LABELS (CONT.)

- **More data stopped helping, and the training recipe was not the reason**
More data kept arriving and the detector kept not improving on it. The training recipe was the first suspect: augmentation, replaying the species the model already knew, completing the missing labels. Each left the same gap, and what the elimination left standing was the data -- the boxes our own auto-labeller writes.
- **A vision model decides whether a dataset belongs, and a chat model cannot**
Professor Zhang's rule: have the system compare what it collected against what we already trust, and discard what sits too far away. The comparison is DINOv2 geometry, not a chat model: Gemma 4 asked to draw the box scores 0.012 F1, and shown a candidate crop it rejects none of the false ones. Fifty-one off-domain datasets were flagged out, the corpus was retrained, and the accuracy did not come back.
- **The unit of curation has to match the unit of noise**
An auto-label error is one wrong box inside an image that is otherwise fine, so a filter reading a whole dataset cannot see it. The per-box tools are written -- crop every box, score the crop against an object-level reference, flag a box whose picture disagrees with its class -- and neither has been re-measured since the bug that made their first run score nothing. That is the open front.
- **A brain filed a falsifiable hypothesis, and it was wrong**
The harvested images carry no plant label: the merge gives every box a class derived from the website it came from. A brain tier proposed that this was what made the curve fall, wrote the condition that would falsify it before any job ran, and filed it for approval. A human approved it and four runs falsified it on that condition. What they exposed instead is density: nine boxes on a harvested image against under two on a clean one.

WHERE THIS GOES: ONE PLATFORM FOR ANY DATASET, ANY FIELD, ANY ROBOT

The end state is one general platform. Agents on it run without being asked, iterate on their own output, escalate to a larger model when they are out of their depth, and are fenced in a sandbox while they do it. Every capability stays in the platform after the experiment that needed it -- the routing table, the ingest contract, the dataset registry, the sandbox, the scorer -- so the next student inherits them working. That is the difference between a lab that has run an experiment and a lab that has built somewhere to run the next one. This is the direction, not a description of what runs today.

Standing today. The collector and trainer are paused after a stop-loss on 2026-08-29. The rover has been off the tailnet since then and the platform answers a drive request with its offline page. The supervisor is advisory. None of the 2,686 frames is labelled.

How the numbers were made. Most in-domain numbers are mAP50-95 on the same 1,977 images under Ultralytics. RF-DETR is pycocotools and the zero-shot column is mAP50 on the 848-image split; neither is differenced against the rest. That holdout is also the validation set during training, so every mAP here is a run's best epoch.

Numbers we withdrew. 0.910 was a data leak, found by a content-level check we built after publishing it. 0.9033 is the best of four unseeded runs and only one of the four crossed 0.90. An earlier reading of the ladder as a flat curve is retracted.