



On one platform, agents collect their own data, review one another's work and train without an operator. The system ran for fifteen unattended rounds; this sheet reports how the agents coordinate and what the collection produced.

Weed Detection



The corpus exceeds 150,000 labelled images. Two agents operate on it unprompted: one harvests and audits imagery, the other trains a detector small enough for the robots.

Rover



A four-wheel skid-steer the width of a crop row carries a Jetson, a Raspberry Pi and its own battery. Camera, GPS and two IMUs stream live over one API key.

Laser Weed Control Robot




A YOLO detector localises each weed; two galvanometer lasers aim and fire, destroying it in place. No chemical is applied and nothing contacts the adjacent crop.

Humanoid and Quadruped



This project pairs a humanoid with a quadruped to record terrain that wheels do not reach. It adopts the Rover's ingest contract, so every session the robots open is governed.

Figure 1. Every project on the platform is created by the same procedure and registers its data under one ingest contract. The harvesting and training agents have so far run on one of them.



Agent Workflow Diagram

Figure 2. A requirement enters. The harvest issues fixed queries without gap analysis. Audited data train the model; the holdout is also the validation set and selects the reported epoch. Evaluation parses a metric; no decision follows. The reviewer runs after every step; the dashed edge is unimplemented.

ONE PLATFORM SUPPORTS ANY ROBOT OR DATASET

- agentAI is one platform for every laboratory robot
- Both the Rover and the laser robot transmit data
- Each drive becomes a dataset at its start
- Every data type and domain forms its own project



The platform is live; the QR code links to it

lab-b660m-c.tailfa6424.ts.net

Authentication requires an institutional Google account.

OPEN-WEIGHT MODELS REVIEW IN TIERS




Figure 3. (e) Recall and (f) grounded recall on one frozen corpus of real incidents from this project's record, scored by its own scorer. Each tail is the language model on artifacts with the twelve checks' findings withheld; each head is the same model supplied with them. The dotted line has no arrowhead: its ends were scored on different cases, the 27 B head on 57 incidents and its tail on 114.

- Review runs even when nothing looks wrong

Review follows every step: failures here announced themselves as successes. Four levels by cost: the running watchdog, twelve fixed checks, a model reading files, then that model given the findings. The watchdog returned nothing on 149 frozen failures.

WHAT THE AGENTS DECIDED ON THEIR OWN

What the lab model says stays a draft	
On the cluster	the answer is authoritative
plan the next step	qwen3.8 27B
review in depth	glm-4.7-flash
hard one-off reasoning	deepseek-v3 671B
decide what to collect	gemma 4
judge dataset quality	qwen2.5 7B
label assistance	minicpm-v
On the lab box	the answer returns marked a draft
turn intent into a project	qwen2.5 3B
summarise a dataset	qwen2.5 3B

Figure 4. The eight jobs that need a model, and the model each one was given, read from the router's own table. A filled square marks a job that exercises judgement. The six on the cluster come back authoritative; the two in the lab come back marked a draft.

- The model's explanation failed its own test

Downloads carry no plant name, so every box from a site gets the same invented label. A model blamed that and wrote down first what would disprove it. A human approved; four runs disproved it. Web pictures average nine boxes, ours under two.

WHAT THE AGENTS DECIDED ON THEIR OWN (CONT.)

- The search ran dry after round three

WHAT THE COLLECTION PRODUCED

- 156,521 labelled by machine, 3,208 kept
- Gemma 4 said yes to every box
- One bad box hides from a dataset filter