



A researcher states a goal; a planning model proposes the project and its agents, and the researcher corrects it. They then run in a sandbox, on data they harvest, take from a robot as it drives, or receive by upload, and train on it round after round unwatched. A loop like that is only as good as what it reports about itself. This sheet is that coordination, and the review that makes it checkable.

Weed Detection



The corpus exists to train the weed detector the laser robot needs. Two agents work it unprompted: one harvests and audits public imagery, the other trains. Over 150,000 labelled images.

Rover



A four-wheel skid-steer narrow enough for a crop row, carrying its own battery and compute. Every drive it opens becomes a governed dataset as it runs: camera, GPS and two inertial units, over one API key.

Laser Weed Control Robot




The end use the collection is for. A YOLO detector localises each weed and two galvanometer lasers aim at it and fire, destroying it in place—no chemical, and nothing touches the crop beside it.

Humanoid and Quadruped



Legs reach ground wheels cannot, and in the field this one crossed it faster than the cart. A wide stance and a level gait hold the camera steady, which is what a collector needs.

Four projects, one platform. Each was opened by the same procedure, and its data is registered the same way whether it was harvested, driven in by a robot, or uploaded.



Agent Workflow Diagram


Figure 1. The harvest-and-train loop, as it runs unattended. A researcher's goal is planned into keyword queries, what they return is audited, what survives is trained on, and the round is scored. The stopping rule is a plateau test today and a reviewer's judgement by design. Review follows every step.

ONE PLATFORM SUPPORTS ANY ROBOT OR DATASET

- agentAI is one platform for every laboratory robot
- The weed detector runs on a robot's live view
- Ask any dataset a question in plain words
- The platform is live; the QR code links to it



OPEN-WEIGHT MODELS REVIEW IN TIERS



WHEN A MODEL'S JUDGEMENT COUNTS

What the lab model says stays a draft	
On the cluster	
plan the next step	qwen3.8 27B
review in depth	glm-4.7-flash
hard one-off reasoning	deepseek-v3 671B
decide what to collect	gemma 4
judge dataset quality	qwen2.5 7B
label assistance	minicpm-v
On the lab box	
turn intent into a project	qwen2.5 3B
summarise a dataset	qwen2.5 3B
The tiers above them	
the hardest open problems	Kimi K3, largest open
check and correct a campaign	Claude Opus 5, via Claude Code
plan the next experiments	Claude Fable 5.1, own key

Figure 3. Each job that needs a model, and the model it was given, from the router's own table; a filled square marks judgement, and the dashed band is the escalation above.

Two jobs run in the lab because a person is waiting; their answers are marked, so none can count as a decision.

What a model may do without asking

Every action carries a risk grade: a reviewer may change a setting, or submit a job if it cites evidence; costly steps wait for approval; deletion is a person's alone. The first to pass was a model's hypothesis, tested and disproved.

WHAT THE COLLECTION PRODUCED

- Two agents collect and train side by side
- Every source is audited before training
- A quality gate decides what the trainer sees
- One bad box hides from a dataset filter