



A researcher states a goal; a planning model proposes the project and its agents, and the researcher corrects it. They then run in a sandbox, on data they harvest, take from a robot as it drives, or receive by upload, and train on it round after round unwatched. A loop like that is only as good as what it reports about itself. This sheet is that coordination, and the review that makes it checkable.

Weed Detection



The corpus exists to train the weed detector the laser robot needs. Two agents work it unprompted: one harvests and audits public imagery, the other trains. Over 150,000 labelled images.

Rover



A four-wheel skid-steer narrow enough for a crop row, carrying its own battery and compute. Every drive it opens becomes a governed dataset as it runs: camera, GPS and two inertial units, over one API key.

Laser Weed Control Robot




The end use the collection is for. A YOLO detector localises each weed and two galvanometer lasers aim at it and fire, destroying it in place—no chemical, and nothing touches the crop beside it.

Humanoid and Quadruped



Legs reach ground wheels cannot, and in the field this one crossed it faster than the cart. A wide stance and a level gait hold the camera steady, which is what a collector needs.

Four projects, one platform. Each was opened by the same procedure, and its data is registered the same way whether it was harvested, driven in by a robot, or uploaded.



Agent Workflow Diagram

Figure 1. The harvest-and-train loop, as it runs unattended. A researcher's goal is planned into keyword queries, what they return is audited, what survives is trained on, and the round is scored. The stopping rule is a plateau test today and a reviewer's judgement by design. Review follows every step.

ONE PLATFORM SUPPORTS ANY ROBOT OR DATASET

agentAI is one platform for every laboratory robot

A robot, drone or sensor rig is a project: its own dataset registry, its own agents, and an owner who controls both. Agents harvest, filter, label and train against that registry. Adding a robot is a configuration change.

The weed detector runs on a robot's live view

While a robot drives, its project page shows the newest camera frame every second and a half. One switch runs the campaign's own weed detector over that frame, so the operator sees what the model sees, on the drive that is happening.

Ask any dataset a question in plain words

Every dataset page has a chat box. A planner reads the question and the data's real columns, then picks a built-in analysis or writes new code. That code runs in a sandbox with no network and no file writes, and the answer returns with its chart and code.



The platform is live; the QR code links to it

lab-b660m-c.tailfa6424.ts.net
Sign in with a Google account to open it.

OPEN-WEIGHT MODELS REVIEW IN TIERS

Citing evidence beats model size




Figure 2. (e) Recall and (f) grounded recall on 149 frozen failures from this project's own record; each arrow runs from a model on raw files to the same model given the checks' findings. A dotted arrow's ends were scored on different cases, 57 against 114.

Handed the twelve checks' findings, a 27-billion model found 0.149 more of the failures and a 14-billion one 0.012.

Review runs even when nothing looks wrong

Review follows every step, because failures here announced themselves as successes. Four levels, by cost: the running watchdog, twelve fixed checks, a model reading raw files, then that model given the checks' findings.

WHEN A MODEL'S JUDGEMENT COUNTS

What the lab model says stays a draft

On the cluster	
plan the next step	qwen3.8 27B
review in depth	glm-4.7-flash
hard one-off reasoning	deepseek-v3 671B
decide what to collect	gemma 4
judge dataset quality	qwen2.5 7B
label assistance	minicpm-v
On the lab box	
turn intent into a project	qwen2.5 3B
summarise a dataset	qwen2.5 3B
The tiers above them	
the hardest open problems	Kimi K3, largest open
check and correct a campaign	Claude Opus 5, via Claude Code
plan the next experiments	Claude Fable 5.1, own key

Figure 3. Each job that needs a model, and the model it was given, from the router's own table; a filled square marks judgement, and the dashed band is the escalation above.

Two jobs run in the lab because a person is waiting; their answers are marked, so none can count as a decision.

What a model may do without asking

Every action carries a risk grade: a reviewer may change a setting, or submit a job if it cites evidence; costly steps wait for approval; deletion is a person's alone. The first to pass was a model's hypothesis, tested and disproved.

WHAT THE COLLECTION PRODUCED

Two agents collect and train side by side

A collector searches Kaggle, Hugging Face, GitHub and Roboflow on its own; a trainer beside it takes each new dataset in between rounds. It has brought in over 600,000 images, and its detection pool holds 179,989 across 63 datasets.

Every image is checked on the way in

Mirrored copies are caught by image content, 65,007 in one pass; no test-set image can reach training; each source is registered with its licence; and boxes are drawn by machine where a dataset had none, 156,521 images so far.

A quality gate decides what the trainer sees

A dataset joins the training pool only if its images resemble trusted weed data in DINOv2 features, which keeps off-topic sets out; fifty-one were held back. A vision model asked the same question said yes to every box.

One bad box hides from a dataset filter

A labelling mistake is usually one bad box in an otherwise sound picture, invisible to a whole-dataset filter. The check has to run box by box. Those tools are built, and what gating the merge on their scores is worth is the next experiment.